

A Static Analysis Approach for Supporting Concurrency in Partitioned Program

Abstract

Program partitioning is an effective way to improve program security. However, there are still many limitations that prevent partitioning framework from being deployed on the majority of modern high-performance programs. In this paper, we identify two key limitations. First, existing partitioning frameworks cannot support concurrent systems due to potential incorrect data synchronization problem. Second, most of partitioning framework normally synchronize more data than necessary, which decrease communication efficiency among partitions and thus cause lower performance. To address these two limitations, we provide two techniques to assist existing partitioning frameworks to support concurrency and avoid synchronizing unneeded data across partition boundary. We implemented our techniques on top of LXDs partitioning framework and evaluate them on a set of Linux drivers. The experiments results shows that our techniques can help existing programs to ...

1. Introduction

Program partitioning refers to the process of separating an existing program into multiple partitions. Each partition runs in a separated domain has its own set of data and privileges. Data are synchronized among different domains through remote procedure call mechanisms. This approach improves program security because compromising one partition doesn't lead to the compromise of other partitions. For programs written in low-level, type-unsafe languages such as C and C++, the partitioning technique is especially beneficial for security as these programs are prone to attack.

There are many attempts on enforcing kernel isolation in recent years. For example, Microdriver... try to isolated drivers from kernel and run the drivers in separate domain.

However, existing kernel isolation frameworks are hard to be deployed correctly on nowadays high-performance operating systems due to many limitations. In this work, we made our observation based on the experiments of previous isolation mechanisms and identify two critical limitations.

First, existing isolation mechanisms do not consider concurrency in their design. In a common isolation architecture, domains synchronize data in a sender-and-receiver analogy: a sender thread and a receiver thread communicate with each other in a serial manner and data synchronization is done synchronously. However, in concurrent programs, asynchronous accesses to data is allowed and thus there is a window of time during which a domain may execute on incorrect data.

Second, most of mechanisms synchronize more data than necessary across partition boundary. When passing large data

structure across partition boundaries, a common synchronization strategy is marshalling/unmarshalling all data structure fields. This greedy approach may incur significant communication overhead when domain transition is frequent in a program. In addition, passing all fields across domain also harms security as high-privileged and sensitive fields within data structures may also be synchronized to low-privileged partitions.

We argue that these two problems need to be addressed in existing isolation mechanisms as they are critical to isolation mechanisms runtime efficiency and correctness.

In this paper, we implemented a tool called SADS, which provides two static analyses to address the two problems we described above. The analyses are implemented in the LXDs kernel isolation framework. We consider our major contribution as follows:

1. Provide a tool called SADS which provides two important static analyses to address data synchronization problems in current kernel isolation frameworks.
2. Analysis 1: computes data that need to be synchronized across domain boundaries during domain communication time. The core of the technique is a static analysis which computes data structure fields that being accessed for each cross-domain function call. One important feature that distinguishes this static analysis from previous works is that it consider domain privacy and prevent synchronizing domain private data at communication time. The key component to allow this is another static analysis for computing domain shared data, which we will describe in details in Section. .
3. Analysis 2: addresses potential data synchronization problems for concurrent programs. In concurrent partitioning programs, data accesses can modified in asynchronous context. This creates potential incorrect data synchronization problem. We illustrate this problem in Section. To address this problem, we provide another static analysis that can identify possible program locations where data synchronization problem can occur and also computes the data that need to be correctly synchronized.

The SADS tool is implemented inside LLVM and integrated with the LXDs kernel partitioning framework. We evaluate the correctness and efficiency of SADS on a set of Linux drivers. The results suggests that SADS can be integrated in practical kernel isolation mechanisms to address potential data synchronization problems.

2. Background and Related Work

In our knowledge, no existing partitioning frameworks have addressed these two problems together.

Program Partitioning PtrSplit is an automatic partitioning system aims that provide marshalling/unmarshalling of pointer type data. The system is capable of generating code that performs marshalling and unmarshalling. While PtrSplit can be applied on many single-threaded program, it fail to address concurrent programs. In addition, the generated communication code by PtrSplit marshalling all fields in data structures that are passed across domain. Thus, none of the two problems are addressed by this system.

Kernel Isolation The Microdriver system [1] provides an architecture for partitioning driver from kernel and allow developers to develop user-mode driver based on it. In this work, a component called code generator is used to generate marshalling/unmarshalling code for data structures passed across partition boundary. This component relies on a static analysis called field accessed analysis, which computes a set of fields that need to be synchronized across domain. Thus, the amount of data structure data passed across boundaries are largely reduced. However, the component has two critical limitations. First, the analysis doesn't consider high-privilege data fields that are private to certain partitions. This indicates that some sensitive data can be synchronized to low-privileged partition. Second, the analysis doesn't describe the computation of fields that may get accessed in concurrent context. Thus, we consider this analysis only partially address the second problem.

3. System Overview

The SADS system provides two core components to address the data synchronization problems in concurrent partitioning program. As the two components are designed and implemented in the form of LLVM pass, SADS can be incorporated into existing partitioning frameworks. Figure.1 shows the overall workflow of SADS. The system first takes the source code of the target program and convert them to LLVM IR using LLVM's frontend. Afterwards, SDAS build a Program Dependency Graph (PDG) using the LLVM IR. Next, two components use the PDG to compute necessary for addressing the data synchronization problems.

The first component is used to improve data communication efficiency among partitions. The component computes a set of necessary fields that need to be passed across partition boundary during inter-domain communications. These data is organized in the form of projection, which is a mechanism for describing data structure fields that need to be synchronized.

The second component addresses potential data synchronization problem for concurrent programs. The component also uses PDG to reason about program locations where data synchronization problems may occur.

The fundamental components for the two static analyses is a

Program Dependency Graph (PDG), which represents the data and control dependencies in the whole program. Based on the PDG, the first component "compute projection" computes data structure fields that need to be synchronized across separation boundary.

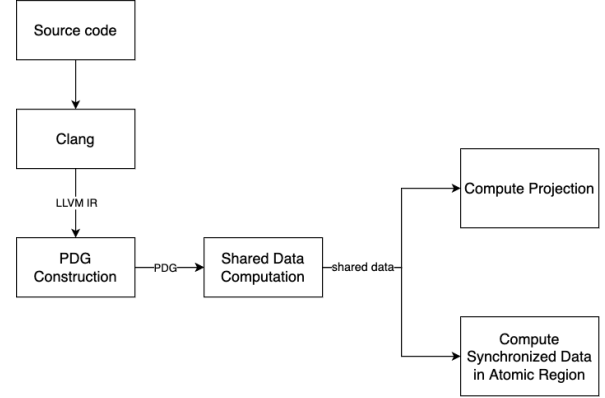


Figure 1: The workflow of SADS toolchain.

4. Design

4.1. PDG Construction

Our analysis for automatic IDL generation is based on the *Program Dependence Graph (PDG)* [?] representation of a program. Conceptually a PDG represents a program's control and data dependence in a single graph. Our PDG construction builds on top of PtrSplit [?] and extends it with field sensitivity, which is necessary for computing projections. Next we provide a brief overview of our PDG construction; details of PDG construction other than field sensitivity can be found in PtrSplit.

Our analysis builds a PDG for an LLVM IR program. In the intraprocedural portion of the PDG, nodes represent IR instructions and edges represent various kinds of control and data dependencies within functions. One example kind is read-after-write data dependence: if instruction n1 reads a piece of memory that was written by n2, there is a data dependence edge from n2 to n1 and the edge is attached with a label that denotes the piece of memory in question.

The interprocedural portion of a PDG adopts an approach called *parameter trees*, a type-based representation that enables modular interprocedural PDG construction. The parameter tree for a parameter is a tree representation of all possible storage locations that a function can access through the parameter. In the example below a struct pointer is passed to function *foo* and the struct has two fields.

```

struct st {int a, float b};
void foo (struct st* s) {...}
  
```

The parameter tree for the parameter *s* is shown in Figure 2. Each node in the parameter tree represents storage that *foo* can

Figure 2: Parameter tree for `foo()`'s parameter `s`.

access through `s`. Therefore, the tree has a node for the pointer to the struct, a node for the struct, and one node for each of the two fields. Parameter trees enable modular interprocedural PDG construction without requiring a global alias analysis, through building actual parameter trees for arguments at a call site and formal parameter trees for parameters of functions (and adding necessary edges); details can be found in [?].

Our static analysis extends the parameter-tree based PDG construction of `PtrSplit` with field sensitivity. In particular, when an instruction accesses a particular field of a struct parameter, a data-dependence edge is added between the instruction and the node that represents the field in the parameter tree. In contrast, being field insensitive, `PtrSplit` adds edges between the instruction and all field nodes in the parameter tree. Field sensitivity is essential for computing projections, which requires an understanding of what fields are read/written by the callee domain when a struct is passed. Our implementation of field sensitivity takes advantage of the fact that, in the LLVM IR, the `GetElementPtr` (GEP) instruction is used to compute the offset of a field; by tracking what offsets are computed by GEP instructions, our analysis resolves the targets of struct reads/writes for most of the cases.

4.2. Identify Cross-domain Function Invocations

The following steps are used to compute a set of cross-domain function calls for each partition:

1. Identify undefined functions in a partition through LLVM.
2. Create a list of library calls.
3. For each undefined function which is not in the list of library calls, we consider it as a cross-domain function call.

On important problem to consider in these steps is finding indirect called functions. For each indirect call, we use the sea-dsa pointer analysis to find out possibly called indirect functions for each indirect call.

4.3. Compute Shared Data

A data is considered shared between two partitions if it has accesses in both of them. The set of shared data specifies what is allowed to be exchanged between two partitions. The role of shared data is essential in both techniques. In projection computation, shared data is used to prevent private data from being synchronized to other partitions. In the process of computing data that need to be synchronized at the end of atomic regions, we only generate summary for shared data.

In general, we use a type-based approach to compute shared data between two partitions. One greedy approach to compute shared data is considering all types of data appear in the program as shared and then computes how this data is accessed. However, this approach is inefficient as it's likely that only a small subset of data types are actually shared. Thus, we first computes a subset of data types that are shared between two

domains, and then computes how the variables of these data types are accessed.

We compute shared data types starting from the possible communication channels between two partitions. In a common partitioned program, two partitions can communicate through two channels: 1. global variables and 2. cross-domain functions. Thus, we consider the data types used in these two channels as possible shared data type. Next, we can further optimize the computation of shared data types by excluding global variables that has uses in only one partition (module). The above steps compute a set of shared data types between two partitions.

After obtaining shared data types, the next step is computing actually shared data. For non-struct types, we simply consider the data is shared between the two domains. For data structure types, we use a global map to store the access information for them; the map key is a data structure type and the value is a set of fields accessed through the type.

The global map is computed through **global type tree**. The algorithm of constructing global type tree is the same as the algorithm of constructing parameter tree. The different is that each global type tree node represents a data field in a data structure type. We mark a node as visited if the corresponding field is accessed in the program.

To compute the accesses to each data field, we add dependency edges between the global tree nodes and variable of the same type. The detailed steps are as follow: We first find a set of IR variables **TV** that match the root node's type in the program and add a **TYPE_MATCH** dependency edge between the root node and the variables. Next, we compute **TV** for the rest of tree node using PDG: for each tree node **n**, we iterate through the user (IR variables) of its parent node's **TV** and find those having the same type and offset with the child node. These variables are included in **n**'s **TV**. Then, we add **TYPE MATCH** dependency edge between a child node and the variable nodes in **TV**. This algorithm is guarantee to terminates since there is not backward edges in global type trees.

4.4. Projections

Compute Projections A field is put into projection if it meets any of the following two conditions:

1. A field is accessed in the transitive closure of a cross-domain function invocation.
2. A field is accessed in asynchronous call context.

The first condition indicates a field is accessed after being passed across domain. Thus, its state is synchronized according to its access type (read/write). The second condition consider possible accesses to shared data in concurrent programs. The high-level steps for computing the projects are as follows:

1. For eah

Optimize Projection Through Shared Data

4.5. Compute Data Need to be Synchronized in Atomic Region

Compute Atomic Region

5. Implementation

5.1.

6. Evaluation

7. Conclusion

8. Acknowledgement

The submission instructions are also available in this [this website](#), including a link to the paper submission site. The website contains sample PDF files for the [paper](#) and [extended abstract](#). The sample files are formatted using the ASPLOS'21 submission format and contain the submission and formatting guidelines. The website also includes an [archive file](#) with \LaTeX templates for both papers and extended abstracts.

All questions regarding paper formatting and submission should be directed to the program co-chairs.

Important highlights:

- Papers should contain a maximum of 10 pages of single-spaced two-column text, including any appendixes, but not including references.
- All submitted papers must be accompanied by an extended abstract, in a separate file with a maximum of 2 pages of single-spaced two-column text, not including references.
- Papers and extended abstracts must be submitted in printable PDF format.
- Text must be in a minimum 10pt (**not** 9pt) font.
- No page limit for references for papers and the extended abstracts.
- Each reference must specify *all* authors (no *et al.*).
- Authors of *all* accepted papers will be required to record a short (less than 2 minute) video that previews the paper. This video substitutes for a lightning talk. Additional requirements for the video will be forthcoming.
- Authors of *all* accepted papers will be required to have a poster in addition to the regular conference talk.
- Proceedings will appear in the ACM digital library up to two weeks before the conference.

Paper evaluation objectives: The committee will make every effort to fairly judge each submitted paper on its own merits. There will be no target acceptance rate. We expect to accept a wide range of papers with appropriate expectations for evaluation. Papers that build on significant past work with strong evaluations are valuable. We encourage you to consider the [SIGPLAN empirical evaluation guidelines](#) for the evaluation of the ideas in your paper. At the same time, papers that open new areas with less rigorous evaluation are equally welcome and especially encouraged. Given the wide range of

topics covered by ASPLOS, every effort will be made to find expert reviewers.

This year, ASPLOS will pilot the use of extended abstracts. **All papers submissions must be accompanied by a 2-page extended abstract, submitted as a separate PDF file.** Extended abstracts will be used throughout the reviewing process so that a larger number of PC members have a better understanding of each paper as they made decisions.

ASPLOS'21 will also feature Artifact Evaluation for accepted papers. Although encouraged, Artifact Evaluation submission is not required nor will it be used as a condition for paper acceptance into ASPLOS 2021. Reviewers will not have visibility into the availability of such artifacts. We request that authors do not refer to them in their paper submissions.

9. Paper and Abstract Preparation Instructions

Formatting instructions and \LaTeX templates for the paper and extended abstract can be found on [this website](#).

9.1. Paper Formatting

Papers must be submitted in printable PDF format and should contain a **maximum of 10 pages** of single-spaced two-column text, including any appendixes, but **not including references**. You may include any number of pages for references, but see below for more instructions. If you are using \LaTeX [2] to typeset your paper, then we suggest that you use [this template](#). If you use a different software package to typeset your paper, then please adhere to the guidelines given in Table 1.

Field	Value
File format	PDF
Page limit	11 pages, not including references
Paper size	US Letter 8.5in \times 11in
Top margin	1in
Bottom margin	1in
Left margin	0.75in
Right margin	0.75in
Body	2-column, single-spaced
Separation between columns	0.25in
Body font	10pt
Abstract font	10pt, italicized
Section heading font	12pt, bold
Subsection heading font	10pt, bold
Caption font	9pt, bold
References	8pt, no page limit, list all authors' names

Table 1: Formatting guidelines for submission.

Please ensure that you include page numbers with your submission. This makes it easier for the reviewers to refer to different parts of your paper when they provide comments.

Please ensure that your submission has a banner at the top of the title page, as shown in [this sample paper](#), which contains the submission number and the notice of confidentiality. If using the template, just replace XXX with your submission number.

9.2. Extended Abstract Formatting

The extended abstracts must be submitted in printable PDF format and should contain a **maximum of 2 pages** of single-spaced two-column text, **not including references**. You may include any number of pages for references, but see below for more instructions. The extended abstracts should use the same formatting as the papers (see [the paper formatting instructions](#)). If you are using L^AT_EX [2] to typeset your extended abstract, then we suggest that you use [this template](#) that also describes that information to include in your extended abstract.

9.3. Content

Author List. Reviewing will be **double blind**; therefore, please **do not include any author names on any submitted documents except in the space provided on the submission form**. You must also ensure that the metadata included in the PDF does not give away the authors. If you are improving upon your prior work, refer to your prior work in the third person and include a full citation for the work in the bibliography. For example, if you are building on *your own* prior work in the papers [3, 4, 5], you would say something like: "While the authors of [3, 4, 5] did X, Y, and Z, this paper additionally does W, and is therefore much better." Do NOT omit or anonymize references for blind review. There is one exception to this for your own prior work that appeared in IEEE CAL, workshops without archived proceedings, etc. as discussed later in this document.

Figures and Tables. Ensure that the figures and tables are legible. Please also ensure that you refer to your figures in the main text. Many reviewers print the papers in gray-scale. Therefore, if you use colors for your figures, ensure that the different colors are highly distinguishable in gray-scale.

References. There is no length limit for references. **Each reference must explicitly list all authors of the paper. Papers not meeting this requirement will be rejected.** Authors of NSF proposals should be familiar with this requirement. Knowing all authors of related work will help find the best reviewers. Since there is no length limit for the number of pages used for references, there is no need to save space here.

10. Paper and Abstract Submission Instructions

10.1. Declaring Authors

Declare all the authors of the paper up front. Addition/removal of authors once the paper is accepted will have to be approved by the program co-chairs, since it potentially undermines the goal of eliminating conflicts for reviewer assignment.

10.2. Areas and Topics

ASPLOS emphasizes multidisciplinary research. Submissions should ideally emphasize synergy of two or more ASPLOS areas: architecture, programming languages, operating systems, and related areas (broadly interpreted). Authors should indicate these areas on the submission form as well as specific topics covered by the paper for optimal reviewer match. If you are unsure whether your paper falls within the scope of ASPLOS, please check with the program co-chair – ASPLOS is a broad, multidisciplinary conference and encourages new topics.

10.3. Declaring Conflicts of Interest

Authors must register all their conflicts on the paper submission site. Conflicts are needed to ensure appropriate assignment of reviewers. If a paper is found to have an undeclared conflict that causes a problem OR if a paper is found to declare false conflicts in order to abuse or “game” the review system, the paper may be rejected.

Please declare a conflict of interest (COI) with the following people for any author of your paper:

1. Your Ph.D. advisor(s), post-doctoral advisor(s), Ph.D. students, and post-doctoral advisees, forever.
2. Family relations by blood or marriage and close friends, forever (if they might be potential reviewers). You are a close friend with someone if you have or would spend a night at their home if you were visiting them, or vice versa.
3. People with whom you have collaborated in the last four years, including
 - co-authors of accepted/rejected/pending papers.
 - co-PIs on accepted/rejected/pending grant proposals.
 - funders (decision-makers) of your research grants, and researchers whom you fund.
4. People (including students) who shared your primary institution(s) in the last four years.

“Service” collaborations such as co-authoring a report for a professional organization, serving on a program committee, or co-presenting tutorials, do not themselves create a conflict of interest. Co-authoring a paper that is a compendium of various projects with no true collaboration among the projects does not constitute a conflict among the authors of the different projects.

On the other hand, there may be others not covered by the above with whom you believe a COI exists, for example, close personal friends. Please report such COIs; however, you may be asked to justify them. Please be reasonable. For example, you cannot declare a COI with a reviewer just because that reviewer works on topics similar to or related to those in your paper. The program co-chairs may contact co-authors to explain a COI whose origin is unclear.

We hope to draw most reviewers from the PC and the ERC, but others from the community may also write reviews. Please declare all your conflicts (not just restricted to the PC and

ERC). When in doubt, contact the program co-chairs.

10.4. Concurrent Submissions and Workshops

By submitting a manuscript to ASPLOS’21, the authors guarantee that the manuscript has not been previously published or accepted for publication in a substantially similar form in any conference, journal, or workshop. The only exceptions are (1) workshops without archived proceedings such as in the ACM digital library (or where the authors chose not to have their paper appear in the archived proceedings), or (2) venues, such as IEEE CAL, where there is an explicit policy that such publication does not preclude longer conference submissions. These are not considered prior publications. Technical reports and papers posted on public social media sites, Web pages, or online repositories, such as arxiv.org, are not considered prior publications either. In these cases, the submitted manuscript may ignore the posted work to preserve author anonymity. The authors also guarantee that no paper that contains significant overlap with the contributions of the submitted paper will be under review for any other conference, journal, or workshop during the ASPLOS’21 review period. Violation of any of these conditions will lead to rejection. As always, if you are in doubt, it is best to contact the program co-chairs. Finally, we also note that the ACM Plagiarism Policy (http://www.acm.org/publications/policies/plagiarism_policy) covers a range of ethical issues concerning the misrepresentation of other works or one’s own work.

10.5. Ethical Obligations

- Authors are not allowed to contact reviewers or PC members to encourage or solicit them to bid on any paper.
- Authors are not allowed to attempt to sway a reviewer to review any paper positively or negatively.
- Authors are not allowed to contact reviewers or PC members requesting any type of information about the reviewing process, either in general or specifically about submitted papers.
- Authors are not allowed to contact reviewers or PC members to ask about the outcomes of any papers.
- Authors must also abide by the **ACM ethics policy**. Violation of the ACM ethics policy may result in rejection of the submission and possible action by the ACM.
- Authors are not allowed to advertise their submissions or related technical reports and postings (e.g., to arxiv.org or online repositories) on social media or community blogs and webpages during the period starting two weeks before the submission deadline and ending when the ASPLOS’21 acceptance results are public.

11. Early Access in the Digital Library

The ASPLOS’21 proceedings will be freely available via the ACM Digital Library for up to two weeks before and up to a month after the conference. **Authors must consider any**

implications of this early disclosure of their work *before* submitting their papers.

12. Acknowledgements

This document is modified from the ASPLOS’20 submission guide, thank you Luis Ceze and Karin Strauss!

References

- [1] Vinod Ganapathy, Matthew J Renzelmann, Arini Balakrishnan, Michael M Swift, and Somesh Jha. The design and implementation of microdrivers. *ACM Sigplan Notices*, 43(3):168–178, 2008.
- [2] Leslie Lamport. *TEX: A Document Preparation System*. Addison-Wesley, Reading, Massachusetts, 2nd edition, 1994.
- [3] Firstname1 Lastname1 and Firstname2 Lastname2. A very nice paper to cite. In *Proceedings of the 33rd Annual ACM SIGPLAN Conference on Programming Language Design and Implementation*, 2012.
- [4] Firstname1 Lastname1, Firstname2 Lastname2, and Firstname3 Lastname3. Another very nice paper to cite. In *Proceedings of the 22nd ACM Symposium on Operating Systems Principles*, 2011.
- [5] Firstname1 Lastname1, Firstname2 Lastname2, Firstname3 Lastname3, Firstname4 Lastname4, and Firstname5 Lastname5. Yet another very nice paper to cite, with many author names all spelled out. In *Proceedings of the 38th Annual International Symposium on Computer Architecture*, 2011.